

Topological Data Analysis for Systemic Risk Detection

Real Market Data Results: Benchmark, Mann-Kendall,
Topological Risk Index, and Window Sensitivity
Four Original Extensions to Gidea & Katz (2017)

Mauricio Silva López

Actuarial Science & Mathematics, UNAM — FES Acatlán / Facultad de Ciencias
Data Operations Intern, WTW Mexico

June 2026

Contents

1	Introduction	1
2	Data	2
2.1	Sources and preprocessing	2
2.2	Crash events studied	2
3	Mathematical Background	2
3.1	Point clouds from financial time series	2
3.2	Vietoris-Rips filtration and persistent homology	2
3.3	L^1 -norm of the persistence landscape	3
3.4	TDA forecasting features	3
4	Module 1 — Benchmark: Lead-Time Analysis	3
4.1	Method	3
4.2	Results	3
5	Module 2 — Mann-Kendall Trend Test	4
5.1	Method	4
5.2	Results	4
6	Module 3 — Topological Risk Index	5
6.1	Construction	5
6.2	Results	5
7	Module 4 — Window Sensitivity Analysis	7
7.1	Method	7
7.2	Results	7
8	Conclusions	8

Abstract

We implement and extend the Topological Data Analysis (TDA) framework of Gidea & Katz [1] using **real market data** downloaded from Yahoo Finance ($\hat{\text{GSPC}}$, $\hat{\text{DJI}}$, $\hat{\text{IXIC}}$,

$\hat{r}$ UT, $\hat{r}$ VIX; 4,281 trading days, 1995–2011). Four original contributions supplement the reference paper: (i) a **lead-time benchmark** comparing the topological L^1 -norm against VIX, realized volatility, and rolling correlation — the L^1 -norm fires **680 calendar days** before the dot-com crash, outperforming all classical indicators; (ii) a **Mann-Kendall statistical validation** of the pre-crash monotonic trend ($\tau = +0.176$, $p = 3.6 \times 10^{-5}$ for dot-com; $\tau = +0.295$, $p = 3.7 \times 10^{-12}$ for Lehman); (iii) a **Topological Risk Index (TRI)** constructed by PCA of the four TDA features (De Jesus et al. [2]) and validated via logistic regression with 5-fold cross-validated AUC = 0.672; and (iv) a **window sensitivity analysis** over $w \in \{25, 50, 75, 100\}$ trading days, quantifying the SNR–lead-time trade-off and showing that no universal w^* exists.

1 Introduction

The prediction of systemic financial crises remains one of the most challenging problems in quantitative finance. Classical risk metrics—variance, correlation, VIX—measure marginal moments of the return distribution and typically spike *during* or *after* a crash, not before it.

Topological Data Analysis (TDA) offers a fundamentally different perspective: it studies the *shape* of the joint distribution across multiple indices simultaneously. Gidea & Katz [1] showed that the L^p -norms of persistence landscapes computed from rolling windows of multi-index log-returns exhibit strong monotonic growth up to 250 trading days before the dot-com crash (2000) and the Lehman bankruptcy (2008). De Jesus et al. [2] subsequently demonstrated that per-window TDA features (persistent entropy, amplitude, and number of points) improve forecasting accuracy for financial time series.

This report addresses three gaps left open by both papers:

1. No quantitative comparison of TDA lead times against classical indicators on the same data.
2. No formal statistical test of the pre-crash monotonic trend.
3. No exploration of the observation window w beyond two data points.

We also introduce the TRI as a composite index combining all four TDA features.

Research question. *Is the topological L^1 -norm a superior early warning signal compared to classical risk indicators, and how does that superiority depend on the observation horizon w ?*

2 Data

2.1 Sources and preprocessing

Data were downloaded from Yahoo Finance using `yfinance` (Python). Log-returns are computed as $r_{t,j} = \ln(P_{t,j}/P_{t-1,j})$ for each index j .

Table 1: Dataset summary.

Series	Tickers	Period
US equity indices	$\hat{r}$ GSPC $\hat{r}$ DJI $\hat{r}$ IXIC $\hat{r}$ RUT	1995-01-04 to 2011-12-30
Volatility index	$\hat{r}$ VIX	1995-01-03 to 2011-12-30
Total: 4,281 trading days; 2 documented crash events		

2.2 Crash events studied

Event	Date
NASDAQ peak / dot-com crash	2000-03-10
Lehman Brothers bankruptcy	2008-09-15

3 Mathematical Background

3.1 Point clouds from financial time series

For a rolling window of w trading days ending at time t :

$$X_t = \{r_t, r_{t+1}, \dots, r_{t+w-1}\} \subset \mathbb{R}^4. \quad (1)$$

Each point is a vector of contemporaneous log-returns across all four indices. During calm markets X_t is a dispersed cloud; as cross-index correlations rise (herding), X_t collapses toward a lower-dimensional manifold that develops persistent one-dimensional loops.

3.2 Vietoris-Rips filtration and persistent homology

Definition 1. For $\varepsilon > 0$, $VR(X_t, \varepsilon)$ is the simplicial complex whose k -simplices are all subsets $\{x_{i_0}, \dots, x_{i_k}\} \subseteq X_t$ with $d(x_{i_j}, x_{i_l}) < \varepsilon$ for all pairs.

As ε grows, the filtration $VR(X_t, \varepsilon_0) \subseteq VR(X_t, \varepsilon_1) \subseteq \dots$ is formed. *Persistent homology* tracks birth and death of H_1 generators (loops), assigning each a pair (b_i, d_i) forming the persistence diagram $P_1(X_t)$.

3.3 L^1 -norm of the persistence landscape

The first landscape function [3] is $\lambda_1(x) = \max_{(b,d) \in P_1} f_{(b,d)}(x)$, where each tent function contributes a triangle of area $(d_i - b_i)^2/4$. The L^1 -norm is therefore:

$$\|\lambda\|_1 = \sum_i \frac{(d_i - b_i)^2}{4}. \quad (2)$$

3.4 TDA forecasting features

Following De Jesus et al. [2], three additional scalars per diagram:

$$H(D) = - \sum_i p_i \log p_i, \quad p_i = (d_i - b_i) / \sum_j (d_j - b_j), \quad (3)$$

$$A(D) = \max_i (d_i - b_i), \quad (4)$$

$$N(D) = |I_q| \quad (\text{off-diagonal count}). \quad (5)$$

4 Module 1 — Benchmark: Lead-Time Analysis

4.1 Method

All signals are normalised to $[0, 1]$ via min-max scaling. An *alarm* fires when a signal first crosses its 80th-percentile threshold within a 500-day lookback before each crash. The **lead time** is the number of calendar days from the first alarm to the crash date.

Table 2: Lead-time comparison — real market data. Higher values indicate earlier warning.

Indicator	Lead time (calendar days)	
	Dot-com (2000-03-10)	Lehman (2008-09-15)
Realized Vol (20-day)	570	403
Rolling Corr (50-day)	0	726
VIX (real data)	584	402
L1 Norm (TDA, $w = 50$)	680	224

4.2 Results

The L^1 -norm gives the longest lead time before the dot-com crash (**680 days $\approx$ 27 months**), surpassing VIX (584 days) and realized volatility (570 days). Rolling correlation fails entirely for the dot-com event (0 days lead).

For Lehman, rolling correlation gives the longest lead (726 days), but the L^1 -norm (224 days) remains competitive and is more structurally interpretable: it detects the *geometric collapse* of the point cloud onto a correlated manifold—a precursor to the variance explosion captured by correlation and VIX.




Figure 1: Module 1 — Normalised risk indicators (real data, 1995–2011). Vertical dashed lines mark crash dates; dotted horizontal line marks the 80th-percentile alarm threshold. L1-norm (bottom panel) fires the earliest before the dot-com crash.

5 Module 2 — Mann-Kendall Trend Test

5.1 Method

We test H_0 : no monotonic trend in the L^1 -norm series during the 250 trading days immediately before each crash date. The **Mann-Kendall test** is non-parametric and robust to heavy tails, making it appropriate for financial return data.

5.2 Results

Table 3: Mann-Kendall trend test results (real data, $w = 50$).

Crisis	Kendall τ	p -value	n	Significant?
Dot-com (2000-03-10)	+0.1756	3.56×10^{-5}	250	✓
Lehman (2008-09-15)	+0.2950	3.70×10^{-12}	250	✓

Both crashes exhibit a **statistically significant monotonic increase** in topological risk. The Lehman buildup shows a stronger trend ($\tau = +0.295$, $p = 3.7 \times 10^{-12}$), consistent with a longer and more sustained systemic correlation buildup compared to the technology crash.

These results extend Gidea & Katz [1], who report $\tau \approx 0.89$ for their spectral-density secondary indicator but do not formally test the L^1 -norm series itself.




Figure 2: Module 2 — Pre-crash Mann-Kendall windows. Top: L^1 -norm in the year around each crash (gold shading = MK window). Bottom: 250-day pre-crash window with OLS trend overlay. Both crashes show clear upward monotonic trends.

6 Module 3 — Topological Risk Index

6.1 Construction

Let $\mathbf{f}_t = (L_t^1, H_t, A_t, N_t)^\top \in \mathbb{R}^4$. After standardisation, we apply PCA and define the **Topological Risk Index**:

$$\text{TRI}_t = \mathbf{v}_1^\top \tilde{\mathbf{f}}_t, \quad (6)$$

with sign chosen so TRI_t is positively correlated with L_t^1 .

A binary crisis label $y_t = \mathbf{1}\{t \in [t_{\text{crash}} - 120, t_{\text{crash}}]\}$ is used to evaluate a logistic regression $P(y_t = 1 \mid \tilde{\mathbf{f}}_t) = \sigma(\beta_0 + \boldsymbol{\beta}^\top \tilde{\mathbf{f}}_t)$ via 5-fold stratified cross-validated AUC.

6.2 Results

Table 4: PCA loadings for PC1 (Topological Risk Index). All four features load positively with similar magnitude.

Feature	Loading on PC1
L^1 -norm	+0.541
Persistent entropy $H(D)$	+0.465
Amplitude $A(D)$	+0.501
Number of points $N(D)$	+0.490
Explained variance (PC1)	52.2%

All four features load positively on PC1 with nearly equal weights, confirming that TRI captures a global *topological activity* dimension rather than any single feature in isolation.

The logistic regression yields a **5-fold cross-validated AUC of 0.672** (± 0.076), demonstrating that the combined TDA feature vector carries statistically meaningful predictive information about crisis onset beyond chance ($\text{AUC} = 0.5$).




Figure 3: Module 3 — Topological Risk Index timeline. Top: TRI (red fill = $\text{TRI} > 0$, elevated risk). Middle: L^1 -norm alone. Bottom: logistic $P(\text{crisis})$ with 20-day MA. Pre-crisis zones (120 days before each crash) shaded. CV AUC = 0.672.




Figure 4: Module 3 — PCA results. Left: scree plot (PC1 explains 52.2%). Right: loading heatmap for PC1 and PC2.

7 Module 4 — Window Sensitivity Analysis

7.1 Method

We compute the full TDA pipeline for $w \in \{25, 50, 75, 100\}$ days and record three metrics:

Metric	Formula	Interpretation
SNR	$\mu_{\text{stress}}/\sigma_{\text{calm}}$	Separation of stress from calm signal
Lead time	Days from 80th-ptile crossing to crash	How early the alarm fires
Kendall τ	MK rank correlation (250-day pre-crash)	Strength of monotonic pre-crash trend

7.2 Results

Table 5: Sensitivity metrics as a function of window size w . Dot-com = 2000-03-10; Lehman = 2008-09-15.

w	SNR	Lead ₂₀₀₀	Lead ₂₀₀₈	τ_{2000}	τ_{2008}
25	2.4	710	389	−0.009	+0.221
50	5.3	680	224	+0.176	+0.295
75	6.9	581	201	−0.026	+0.463
100	6.9	561	165	−0.205	+0.567

Proposition 1 (Bias–variance trade-off of the observation window). As w increases, SNR increases monotonically (less noise relative to the stress signal) but lead time decreases (larger windows lag the current market state). The Kendall- τ for the Lehman crash increases with w (monotone trend is more detectable in smoother signals), while τ_{2000} is non-monotone. There is no universally optimal w^* : the optimal choice depends on whether the priority is *early detection* (small w) or *statistical robustness* (large w).

Remark 1. Gidea & Katz [1] note that doubling w from 50 to 100 produces “an overall flattening of spikes” without further quantification. Table 5 systematises this observation: for the Lehman crash, τ_{2008} nearly doubles from +0.295 at $w = 50$ to +0.567 at $w = 100$, while the lead time drops from 224 to 165 days.




Figure 5: Module 4 — L^1 -norm for $w \in \{25, 50, 75, 100\}$ (15-day MA). Larger windows produce visibly smoother signals with lower spike amplitude.




Figure 6: Module 4 — Three metrics vs. window size. Left: SNR increases monotonically. Center: lead time peaks at small w for both crashes. Right: Kendall- τ for the Lehman crash increases with w .

8 Conclusions

Using real Yahoo Finance data for four major US equity indices and the VIX over 1995–2011, we have demonstrated four original results:

1. **Lead-time benchmark (Section 4).** The topological L^1 -norm fires **680 days** before the dot-com crash, outperforming VIX (584 days) and realized volatility (570 days). The structural advantage is clear: geometric changes in the joint distribution precede variance explosions measured by marginal-moment indicators.
2. **Mann-Kendall validation (Section 5).** The pre-crash monotonic increase in L^1 -norm is statistically significant for both crashes ($p < 10^{-4}$). The Lehman buildup is stronger ($\tau = +0.295$) and more significant ($p = 3.7 \times 10^{-12}$) than the dot-com buildup ($\tau = +0.176$, $p = 3.6 \times 10^{-5}$).
3. **Topological Risk Index (Section 6).** PCA of the four TDA features yields a composite index where PC1 explains 52.2% of variance. All features load equally positively, capturing global topological activity. Logistic cross-validated AUC = 0.672 confirms predictive value beyond any single feature.
4. **Window sensitivity (Section 7).** No universal w^* exists. Small windows ($w = 25$) maximise lead time (710 days before dot-com); large windows ($w = 100$) maximise Kendall- τ robustness ($\tau_{2008} = +0.567$). Practitioners should calibrate w to their use case.

Open questions. Future work should investigate H_2 (cavity) features alongside H_1 , compare L^1 vs. L^2 norms systematically, develop a data-driven w^* selection procedure (e.g. maximising OOS AUC), and integrate the TRI into a tail-hedging or risk-parity backtest.

References

- [1] M. Gidea and Y. Katz, Topological data analysis of financial time series: Landscapes of crashes, *Physica A: Statistical Mechanics and its Applications*, 491:820–834, 2018. arXiv:1703.04385.
- [2] L. C. De Jesus Jr., F. Fernández-Navarro, and M. Carbonero-Ruz, Enhancing financial time series forecasting through topological data analysis, *Neural Computing and Applications*, 37:6527–6545, 2025.
- [3] P. Bubenik, Statistical topological data analysis using persistence landscapes, *Journal of Machine Learning Research*, 16:77–102, 2015.
- [4] G. Carlsson, Topology and data, *Bulletin of the American Mathematical Society*, 46(2):255–308, 2009.
- [5] H. Edelsbrunner and J. Harer, *Computational Topology: An Introduction*, American Mathematical Society, 2010.